

# Big Data Project (21L-5625,L21-5626,L21-5632)

## Movie and Political Sentiment Analysis with PySpark

```
%%bash
#Installing Pyspark
pip install pyspark

Looking in indexes: https://pypi.org/simple, https://us-
python.pkg.dev/colab-wheels/public/simple/
Collecting pyspark
  Downloading pyspark-3.4.0.tar.gz (310.8 MB)
     310.8/310.8 MB 2.9 MB/s
eta 0:00:00
  Preparing metadata (setup.py): started
  Preparing metadata (setup.py): finished with status 'done'
Requirement already satisfied: py4j==0.10.9.7 in
/usr/local/lib/python3.9/dist-packages (from pyspark) (0.10.9.7)
Building wheels for collected packages: pyspark
  Building wheel for pyspark (setup.py): started
  Building wheel for pyspark (setup.py): finished with status 'done'
  Created wheel for pyspark: filename=pyspark-3.4.0-py2.py3-none-
any.whl size=311317145
sha256=83ae4052984f025147841752d80310a153888e6e34a3a4e1540182459372538
7
  Stored in directory:
/root/.cache/pip/wheels/9f/34/a4/159aa12d0a510d5ff7c8f0220abbea42e5d81
ecf588c4fd884
Successfully built pyspark
Installing collected packages: pyspark
Successfully installed pyspark-3.4.0

%%bash

# Download the data files from github
#or one can just run this python notebook, since I am just linking
the sources
data_file_1=imdb_reviews_preprocessed.parquet
data_file_2=sentiments.parquet
data_file_3=tweets.parquet

# If data_file_1 file does not exist in the colab environment
if [[ ! -f ${data_file_1} ]]; then
  # download the data file from github and save it in this colab
environment instance
  wget
  https://raw.githubusercontent.com/wewilllil/ist718_data/master/${
data_file_1}
```

```

fi

# If data_file_1 file does not exist in the colab environment
if [[ ! -f ${data_file_2} ]]; then
    # download the data file from github and save it in this colab
    # environment instance
    wget
    https://raw.githubusercontent.com/wewilllil/ist718_data/master/${data_file_2}
fi

# If data_file_1 file does not exist in the colab environment
if [[ ! -f ${data_file_3} ]]; then
    # download the data file from github and save it in this colab
    # environment instance
    wget
    https://raw.githubusercontent.com/wewilllil/ist718_data/master/${data_file_3}
fi

--2023-04-24 19:19:31--
https://raw.githubusercontent.com/wewilllil/ist718_data/master/imdb_reviews_preprocessed.parquet
Resolving raw.githubusercontent.com (raw.githubusercontent.com)...
185.199.108.133, 185.199.109.133, 185.199.110.133, ...
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|
185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 21597134 (21M) [application/octet-stream]
Saving to: 'imdb_reviews_preprocessed.parquet'

 0K ..... 0%
8.57M 2s
 50K ..... 0%
8.00M 2s
100K ..... 0%
42.3M 2s
150K ..... 0%
45.5M 1s
200K ..... 1%
12.5M 1s
250K ..... 1%
52.1M 1s
300K ..... 1%
47.8M 1s
350K ..... 1%
89.3M 1s
400K ..... 2%
233M 1s
450K ..... 2%

```

14.3M 1s	
500K .....	2%
50.0M 1s	
550K .....	2%
62.2M 1s	
600K .....	3%
199M 1s	
650K .....	3%
204M 1s	
700K .....	3%
235M 1s	
750K .....	3%
220M 1s	
800K .....	4%
231M 1s	
850K .....	4%
191M 1s	
900K .....	4%
21.4M 1s	
950K .....	4%
47.6M 1s	
1000K .....	4%
38.9M 1s	
1050K .....	5%
41.5M 1s	
1100K .....	5%
18.2M 1s	
1150K .....	5%
40.4M 1s	
1200K .....	5%
51.0M 1s	
1250K .....	6%
33.3M 1s	
1300K .....	6%
47.4M 1s	
1350K .....	6%
42.8M 1s	
1400K .....	6%
37.6M 1s	
1450K .....	7%
46.8M 1s	
1500K .....	7%
138M 1s	
1550K .....	7%
163M 1s	
1600K .....	7%
214M 1s	
1650K .....	8%
184M 1s	

1700K	8%
84.1M 1s	
1750K	8%
44.6M 0s	
1800K	8%
50.9M 0s	
1850K	9%
36.7M 0s	
1900K	9%
59.6M 0s	
1950K	9%
42.2M 0s	
2000K	9%
53.8M 0s	
2050K	9%
53.7M 0s	
2100K	10%
50.0M 0s	
2150K	10%
48.4M 0s	
2200K	10%
219M 0s	
2250K	10%
198M 0s	
2300K	11%
191M 0s	
2350K	11%
245M 0s	
2400K	11%
248M 0s	
2450K	11%
194M 0s	
2500K	12%
179M 0s	
2550K	12%
73.6M 0s	
2600K	12%
47.3M 0s	
2650K	12%
52.5M 0s	
2700K	13%
45.1M 0s	
2750K	13%
52.4M 0s	
2800K	13%
51.9M 0s	
2850K	13%
47.3M 0s	
2900K	13%

51.2M 0s	
2950K .....	14%
46.3M 0s	
3000K .....	14%
52.6M 0s	
3050K .....	14%
65.9M 0s	
3100K .....	14%
239M 0s	
3150K .....	15%
214M 0s	
3200K .....	15%
184M 0s	
3250K .....	15%
249M 0s	
3300K .....	15%
240M 0s	
3350K .....	16%
209M 0s	
3400K .....	16%
177M 0s	
3450K .....	16%
244M 0s	
3500K .....	16%
243M 0s	
3550K .....	17%
197M 0s	
3600K .....	17%
262M 0s	
3650K .....	17%
147M 0s	
3700K .....	17%
239M 0s	
3750K .....	18%
248M 0s	
3800K .....	18%
242M 0s	
3850K .....	18%
177M 0s	
3900K .....	18%
203M 0s	
3950K .....	18%
247M 0s	
4000K .....	19%
248M 0s	
4050K .....	19%
91.3M 0s	
4100K .....	19%
49.0M 0s	

4150K .....	19%
55.6M 0s	
4200K .....	20%
46.8M 0s	
4250K .....	20%
48.6M 0s	
4300K .....	20%
56.5M 0s	
4350K .....	20%
43.8M 0s	
4400K .....	21%
51.1M 0s	
4450K .....	21%
52.3M 0s	
4500K .....	21%
55.3M 0s	
4550K .....	21%
52.3M 0s	
4600K .....	22%
43.3M 0s	
4650K .....	22%
52.5M 0s	
4700K .....	22%
50.4M 0s	
4750K .....	22%
52.8M 0s	
4800K .....	22%
53.6M 0s	
4850K .....	23%
47.3M 0s	
4900K .....	23%
51.4M 0s	
4950K .....	23%
49.4M 0s	
5000K .....	23%
52.6M 0s	
5050K .....	24%
55.3M 0s	
5100K .....	24%
47.6M 0s	
5150K .....	24%
57.6M 0s	
5200K .....	24%
128M 0s	
5250K .....	25%
188M 0s	
5300K .....	25%
250M 0s	
5350K .....	25%

223M 0s	
5400K .....	25%
243M 0s	
5450K .....	26%
233M 0s	
5500K .....	26%
170M 0s	
5550K .....	26%
224M 0s	
5600K .....	26%
253M 0s	
5650K .....	27%
223M 0s	
5700K .....	27%
215M 0s	
5750K .....	27%
203M 0s	
5800K .....	27%
190M 0s	
5850K .....	27%
250M 0s	
5900K .....	28%
222M 0s	
5950K .....	28%
239M 0s	
6000K .....	28%
201M 0s	
6050K .....	28%
253M 0s	
6100K .....	29%
245M 0s	
6150K .....	29%
102M 0s	
6200K .....	29%
34.4M 0s	
6250K .....	29%
51.1M 0s	
6300K .....	30%
48.2M 0s	
6350K .....	30%
90.8M 0s	
6400K .....	30%
131M 0s	
6450K .....	30%
67.4M 0s	
6500K .....	31%
49.1M 0s	
6550K .....	31%
53.3M 0s	

6600K .....	31%
44.3M 0s	
6650K .....	31%
68.3M 0s	
6700K .....	32%
177M 0s	
6750K .....	32%
165M 0s	
6800K .....	32%
197M 0s	
6850K .....	32%
50.8M 0s	
6900K .....	32%
55.7M 0s	
6950K .....	33%
54.3M 0s	
7000K .....	33%
40.8M 0s	
7050K .....	33%
93.2M 0s	
7100K .....	33%
93.3M 0s	
7150K .....	34%
55.1M 0s	
7200K .....	34%
43.8M 0s	
7250K .....	34%
56.3M 0s	
7300K .....	34%
56.9M 0s	
7350K .....	35%
59.6M 0s	
7400K .....	35%
65.3M 0s	
7450K .....	35%
82.9M 0s	
7500K .....	35%
75.9M 0s	
7550K .....	36%
73.7M 0s	
7600K .....	36%
60.7M 0s	
7650K .....	36%
5.65M 0s	
7700K .....	36%
60.9M 0s	
7750K .....	36%
59.6M 0s	
7800K .....	37%



54.2M 0s	
7850K .....	37%
46.9M 0s	
7900K .....	37%
58.7M 0s	
7950K .....	37%
23.8M 0s	
8000K .....	38%
52.2M 0s	
8050K .....	38%
50.8M 0s	
8100K .....	38%
119M 0s	
8150K .....	38%
179M 0s	
8200K .....	39%
216M 0s	
8250K .....	39%
203M 0s	
8300K .....	39%
199M 0s	
8350K .....	39%
192M 0s	
8400K .....	40%
224M 0s	
8450K .....	40%
218M 0s	
8500K .....	40%
130M 0s	
8550K .....	40%
119M 0s	
8600K .....	41%
137M 0s	
8650K .....	41%
252M 0s	
8700K .....	41%
256M 0s	
8750K .....	41%
253M 0s	
8800K .....	41%
216M 0s	
8850K .....	42%
227M 0s	
8900K .....	42%
269M 0s	
8950K .....	42%
233M 0s	
9000K .....	42%
213M 0s	

9050K	43%
244M 0s	
9100K	43%
195M 0s	
9150K	43%
259M 0s	
9200K	43%
246M 0s	
9250K	44%
237M 0s	
9300K	44%
182M 0s	
9350K	44%
231M 0s	
9400K	44%
252M 0s	
9450K	45%
251M 0s	
9500K	45%
245M 0s	
9550K	45%
72.9M 0s	
9600K	45%
55.2M 0s	
9650K	45%
40.3M 0s	
9700K	46%
52.9M 0s	
9750K	46%
37.2M 0s	
9800K	46%
50.8M 0s	
9850K	46%
49.1M 0s	
9900K	47%
9.68M 0s	
9950K	47%
53.8M 0s	
10000K	47%
47.8M 0s	
10050K	47%
54.7M 0s	
10100K	48%
44.7M 0s	
10150K	48%
51.5M 0s	
10200K	48%
39.9M 0s	
10250K	48%

55.6M 0s	
10300K .....	49%
55.1M 0s	
10350K .....	49%
5.45M 0s	
10400K .....	49%
47.6M 0s	
10450K .....	49%
57.1M 0s	
10500K .....	50%
55.9M 0s	
10550K .....	50%
55.7M 0s	
10600K .....	50%
49.5M 0s	
10650K .....	50%
65.3M 0s	
10700K .....	50%
65.5M 0s	
10750K .....	51%
50.9M 0s	
10800K .....	51%
61.3M 0s	
10850K .....	51%
80.1M 0s	
10900K .....	51%
93.1M 0s	
10950K .....	52%
237M 0s	
11000K .....	52%
197M 0s	
11050K .....	52%
250M 0s	
11100K .....	52%
215M 0s	
11150K .....	53%
253M 0s	
11200K .....	53%
249M 0s	
11250K .....	53%
77.8M 0s	
11300K .....	53%
47.5M 0s	
11350K .....	54%
45.0M 0s	
11400K .....	54%
54.7M 0s	
11450K .....	54%
54.8M 0s	

11500K .....	54%
54.5M 0s	
11550K .....	54%
48.5M 0s	
11600K .....	55%
43.1M 0s	
11650K .....	55%
55.7M 0s	
11700K .....	55%
152M 0s	
11750K .....	55%
224M 0s	
11800K .....	56%
193M 0s	
11850K .....	56%
263M 0s	
11900K .....	56%
271M 0s	
11950K .....	56%
213M 0s	
12000K .....	57%
246M 0s	
12050K .....	57%
230M 0s	
12100K .....	57%
264M 0s	
12150K .....	57%
227M 0s	
12200K .....	58%
255M 0s	
12250K .....	58%
241M 0s	
12300K .....	58%
216M 0s	
12350K .....	58%
216M 0s	
12400K .....	59%
262M 0s	
12450K .....	59%
257M 0s	
12500K .....	59%
248M 0s	
12550K .....	59%
200M 0s	
12600K .....	59%
261M 0s	
12650K .....	60%
252M 0s	
12700K .....	60%
261M 0s	

12750K	60%
86.0M 0s	
12800K	60%
41.6M 0s	
12850K	61%
53.6M 0s	
12900K	61%
75.1M 0s	
12950K	61%
209M 0s	
13000K	61%
169M 0s	
13050K	62%
184M 0s	
13100K	62%
195M 0s	
13150K	62%
222M 0s	
13200K	62%
219M 0s	
13250K	63%
187M 0s	
13300K	63%
66.6M 0s	
13350K	63%
35.4M 0s	
13400K	63%
47.8M 0s	
13450K	64%
52.4M 0s	
13500K	64%
52.7M 0s	
13550K	64%
56.0M 0s	
13600K	64%
42.5M 0s	
13650K	64%
47.0M 0s	
13700K	65%
21.0M 0s	
13750K	65%
55.6M 0s	
13800K	65%
57.7M 0s	
13850K	65%
38.8M 0s	
13900K	66%
53.6M 0s	
13950K	66%

55.3M 0s	
14000K .....	66%
50.2M 0s	
14050K .....	66%
66.6M 0s	
14100K .....	67%
55.7M 0s	
14150K .....	67%
83.5M 0s	
14200K .....	67%
130M 0s	
14250K .....	67%
205M 0s	
14300K .....	68%
244M 0s	
14350K .....	68%
244M 0s	
14400K .....	68%
264M 0s	
14450K .....	68%
218M 0s	
14500K .....	68%
237M 0s	
14550K .....	69%
242M 0s	
14600K .....	69%
224M 0s	
14650K .....	69%
264M 0s	
14700K .....	69%
216M 0s	
14750K .....	70%
258M 0s	
14800K .....	70%
194M 0s	
14850K .....	70%
234M 0s	
14900K .....	70%
164M 0s	
14950K .....	71%
256M 0s	
15000K .....	71%
252M 0s	
15050K .....	71%
191M 0s	
15100K .....	71%
249M 0s	
15150K .....	72%
233M 0s	

15200K	72%
255M 0s	
15250K	72%
254M 0s	
15300K	72%
48.5M 0s	
15350K	73%
44.1M 0s	
15400K	73%
49.0M 0s	
15450K	73%
50.2M 0s	
15500K	73%
48.4M 0s	
15550K	73%
55.1M 0s	
15600K	74%
55.0M 0s	
15650K	74%
39.8M 0s	
15700K	74%
59.1M 0s	
15750K	74%
52.5M 0s	
15800K	75%
45.2M 0s	
15850K	75%
54.6M 0s	
15900K	75%
49.4M 0s	
15950K	75%
47.4M 0s	
16000K	76%
55.5M 0s	
16050K	76%
54.0M 0s	
16100K	76%
41.9M 0s	
16150K	76%
53.3M 0s	
16200K	77%
55.1M 0s	
16250K	77%
50.7M 0s	
16300K	77%
49.6M 0s	
16350K	77%
54.6M 0s	
16400K	77%

56.8M 0s	
16450K .....	78%
78.1M 0s	
16500K .....	78%
182M 0s	
16550K .....	78%
249M 0s	
16600K .....	78%
240M 0s	
16650K .....	79%
249M 0s	
16700K .....	79%
224M 0s	
16750K .....	79%
226M 0s	
16800K .....	79%
211M 0s	
16850K .....	80%
243M 0s	
16900K .....	80%
246M 0s	
16950K .....	80%
214M 0s	
17000K .....	80%
220M 0s	
17050K .....	81%
250M 0s	
17100K .....	81%
215M 0s	
17150K .....	81%
238M 0s	
17200K .....	81%
179M 0s	
17250K .....	82%
260M 0s	
17300K .....	82%
210M 0s	
17350K .....	82%
243M 0s	
17400K .....	82%
192M 0s	
17450K .....	82%
55.3M 0s	
17500K .....	83%
56.7M 0s	
17550K .....	83%
49.3M 0s	
17600K .....	83%
56.0M 0s	



17650K .....	83%
56.9M 0s	
17700K .....	84%
56.7M 0s	
17750K .....	84%
54.2M 0s	
17800K .....	84%
49.5M 0s	
17850K .....	84%
55.4M 0s	
17900K .....	85%
56.0M 0s	
17950K .....	85%
55.2M 0s	
18000K .....	85%
55.9M 0s	
18050K .....	85%
48.8M 0s	
18100K .....	86%
56.5M 0s	
18150K .....	86%
52.1M 0s	
18200K .....	86%
54.5M 0s	
18250K .....	86%
56.8M 0s	
18300K .....	87%
48.7M 0s	
18350K .....	87%
54.9M 0s	
18400K .....	87%
79.1M 0s	
18450K .....	87%
208M 0s	
18500K .....	87%
211M 0s	
18550K .....	88%
245M 0s	
18600K .....	88%
238M 0s	
18650K .....	88%
241M 0s	
18700K .....	88%
180M 0s	
18750K .....	89%
250M 0s	
18800K .....	89%
230M 0s	
18850K .....	89%

245M 0s	
18900K .....	89%
236M 0s	
18950K .....	90%
160M 0s	
19000K .....	90%
249M 0s	
19050K .....	90%
246M 0s	
19100K .....	90%
240M 0s	
19150K .....	91%
199M 0s	
19200K .....	91%
239M 0s	
19250K .....	91%
254M 0s	
19300K .....	91%
250M 0s	
19350K .....	91%
204M 0s	
19400K .....	92%
54.0M 0s	
19450K .....	92%
56.4M 0s	
19500K .....	92%
56.2M 0s	
19550K .....	92%
49.5M 0s	
19600K .....	93%
56.6M 0s	
19650K .....	93%
56.3M 0s	
19700K .....	93%
49.1M 0s	
19750K .....	93%
53.6M 0s	
19800K .....	94%
57.0M 0s	
19850K .....	94%
55.3M 0s	
19900K .....	94%
49.1M 0s	
19950K .....	94%
56.1M 0s	
20000K .....	95%
55.7M 0s	
20050K .....	95%
56.8M 0s	

20100K	95%
48.8M 0s	
20150K	95%
56.9M 0s	
20200K	96%
56.0M 0s	
20250K	96%
52.3M 0s	
20300K	96%
49.4M 0s	
20350K	96%
75.6M 0s	
20400K	96%
221M 0s	
20450K	97%
251M 0s	
20500K	97%
202M 0s	
20550K	97%
251M 0s	
20600K	97%
248M 0s	
20650K	98%
196M 0s	
20700K	98%
243M 0s	
20750K	98%
240M 0s	
20800K	98%
252M 0s	
20850K	99%
216M 0s	
20900K	99%
217M 0s	
20950K	99%
248M 0s	
21000K	99%
242M 0s	
21050K	100%
252M=0.3s	

2023-04-24 19:19:33 (70.1 MB/s) - 'imdb\_reviews\_preprocessed.parquet'  
saved [21597134/21597134]

--2023-04-24 19:19:33--

[https://raw.githubusercontent.com/wewilllil/ist718\\_data/master/sentiments.parquet](https://raw.githubusercontent.com/wewilllil/ist718_data/master/sentiments.parquet)

Resolving raw.githubusercontent.com (raw.githubusercontent.com)...  
185.199.108.133, 185.199.109.133, 185.199.110.133, ...

```
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|
185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 56439 (55K) [application/octet-stream]
Saving to: 'sentiments.parquet'
```

```
0K ..... 90%
5.71M 0s
50K ..... 100%
5.29M=0.009s
```

```
2023-04-24 19:19:33 (5.67 MB/s) - 'sentiments.parquet' saved
[56439/56439]
```

```
--2023-04-24 19:19:33--
```

```
https://raw.githubusercontent.com/wewillil/ist718_data/master/tweets.p
arquet
```

```
Resolving raw.githubusercontent.com (raw.githubusercontent.com)...
185.199.108.133, 185.199.109.133, 185.199.110.133, ...
```

```
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|
185.199.108.133|:443... connected.
```

```
HTTP request sent, awaiting response... 200 OK
```

```
Length: 154653 (151K) [application/octet-stream]
```

```
Saving to: 'tweets.parquet'
```

```
0K ..... 33%
6.55M 0s
50K ..... 66%
6.82M 0s
100K ..... 99%
40.0M 0s
150K . 100%
1961G=0.02s
```

```
2023-04-24 19:19:33 (9.31 MB/s) - 'tweets.parquet' saved
[154653/154653]
```

## Loading packages and connecting to Spark cluster

```
from __future__ import division
from pyspark.sql import SparkSession
from pyspark.ml import feature, regression, evaluation, Pipeline
from pyspark.sql import functions as fn, Row
import matplotlib.pyplot as plt
import glob
import subprocess
import numpy as np
import pandas as pd
import os
```

```
spark = SparkSession.builder.getOrCreate()
sc = spark.sparkContext
```

## Load sentiment data

```
sentiments_df = spark.read.parquet('sentiments.parquet')
sentiments_df.printSchema()
```

```
root
|-- word: string (nullable = true)
|-- sentiment: long (nullable = true)
```

```
sentiments_df.show(5)
```

```
+-----+-----+
|      word|sentiment|
+-----+-----+
| gratefully|         1|
| gratification|         1|
| gratified|         1|
| gratifies|         1|
| gratify|         1|
+-----+-----+
```

only showing top 5 rows

The schema is very simple: for each word, we have whether it is positive (+1) or negative (-1)

```
# a sample of positive words
```

```
sentiments_df.where(fn.col('sentiment') == 1).show(5)
```

```
+-----+-----+
|      word|sentiment|
+-----+-----+
| gratefully|         1|
| gratification|         1|
| gratified|         1|
| gratifies|         1|
| gratify|         1|
+-----+-----+
```

only showing top 5 rows

```
# a sample of negative words
```

```
sentiments_df.where(fn.col('sentiment') == -1).show(5)
```

```
+-----+-----+
|      word|sentiment|
```

```

+-----+-----+
| 2-faced |      -1 |
| 2-faces |      -1 |
| abnormal |      -1 |
| abolish |      -1 |
| abominable |    -1 |
+-----+-----+
only showing top 5 rows

```

#Lets see how many of each category we have

```
sentiments_df.groupBy('sentiment').agg(fn.count('*')).show()
```

```

+-----+-----+
| sentiment | count(1) |
+-----+-----+
|          1 |      2006 |
|         -1 |      4783 |
+-----+-----+

```

We have almost two times the number of negative words!

To test our approach, we will use a sample of [IMDB](#) reviews that were tagged as positive and negative.

Let's load them:

```

imdb_reviews_df =
spark.read.parquet('imdb_reviews_preprocessed.parquet')
imdb_reviews_df.show(10)

```

```

+-----+-----+-----+
|      id |      review | score |
+-----+-----+-----+
| pos_10006 | In this "critical... | 1.0 |
| pos_10013 | Like one of the p... | 1.0 |
| pos_10022 | Aro Tolbukhin bur... | 1.0 |
| pos_10033 | The movie Titanic... | 1.0 |
| pos_1003 | Another Aussie ma... | 1.0 |
| pos_1004 | After a brief pro... | 1.0 |
| pos_10053 | I must admit, whe... | 1.0 |
| pos_10062 | Wow. What a wonde... | 1.0 |
| pos_10074 | quote by Nicolas ... | 1.0 |
| pos_10083 | The fact that thi... | 1.0 |
+-----+-----+-----+
only showing top 10 rows

```

The schema is very simple: for each word, we have whether it is positive (+1) or negative (-1)

Print the unique scores in the `imdb_reviews_df`. A positive review has a score of 1, and a negative review has a score of 0.

```
imdb_reviews_df.toPandas()['score'].unique()  
array([1., 0.]
```

Let's take a look at a positive review

```
imdb_reviews_df.where(fn.col('score') == 1).first()  
Row(id='pos_10006', review='In this "critically acclaimed  
psychological thriller based on true events, Gabriel (Robin Williams),  
a celebrated writer and late-night talk show host, becomes captivated  
by the harrowing story of a young listener and his adoptive mother  
(Toni Collette). When troubling questions arise about this boy\'s  
(story), however, Gabriel finds himself drawn into a widening mystery  
that hides a deadly secret\x85" according to film\'s official  
synopsis.<br /><br />You really should STOP reading these comments,  
and watch the film NOW...<br /><br />The "How did he lose his leg?"  
ending, with Ms. Collette planning her new life, should be chopped  
off, and sent to "deleted scenes" land. It\'s overkill. The true  
nature of her physical and mental ailments should be obvious, by the  
time Mr. Williams returns to New York. Possibly, her blindness could  
be in question - but a revelation could have been made certain in either  
the "highway" or "video tape" scenes. The film would benefit from a  
re-editing - how about a "director\'s cut"? <br /><br />Williams and  
Bobby Cannavale (as Jess) don\'t seem, initially, believable as a  
couple. A scene or two establishing their relationship might have  
helped set the stage. Otherwise, the cast is exemplary. Williams  
offers an exceptionally strong characterization, and not a "gay  
impersonation". Sandra Oh (as Anna), Joe Morton (as Ashe), and Rory  
Culkin (Pete Logand) are all perfect.<br /><br />Best of all,  
Collette\'s "Donna" belongs in the creepy hall of fame. Ms. Oh is  
correct in saying Collette might be, "you know, like that guy  
from \'Psycho\'. There have been several years when organizations  
giving acting awards seemed to reach for women, due to a slighter  
dispersion of roles; certainly, they could have noticed Collette with  
some award consideration. She is that good. And, director Patrick  
Stettner definitely evokes Hitchcock - he even makes getting a  
sandwich from a vending machine suspenseful.<br /><br />Finally,  
writers Stettner, Armistead Maupin, and Terry Anderson deserve  
gratitude from flight attendants everywhere.<br /><br />***** The  
Night Listener (1/21/06) Patrick Stettner ~ Robin Williams, Toni  
Collette, Sandra Oh, Rory Culkin', score=1.0)
```

And a negative one

```
imdb_reviews_df.where(fn.col('score') == 0).first()
```

```
Row(id='neg_10006', review="I don't know who to blame, the timid writers or the clueless director. It seemed to be one of those movies where so much was paid to the stars (Angie, Charlie, Denise, Rosanna and Jon) that there wasn't enough left to really make a movie. This could have been very entertaining, but there was a veil of timidity, even cowardice, that hung over each scene. Since it got an R rating anyway why was the ubiquitous bubble bath scene shot with a 70-year-old woman and not Angie Harmon? Why does Sheen sleepwalk through potentially hot relationships WITH TWO OF THE MOST BEAUTIFUL AND SEXY ACTRESSES in the world? If they were only looking for laughs why not cast Whoopi Goldberg and Judy Tenuta instead? This was so predictable I was surprised to find that the director wasn't a five year old. What a waste, not just for the viewers but for the actors as well.", score=0.0)
```

The first problem that we encounter is that the reviews are in plain text. We need to split the words and then match them to `sentiment_df`. To do this, we will use a transformation that takes raw text and outputs a list of words

```
from pyspark.ml.feature import RegexTokenizer
```

`RegexTokenizer` extracts a sequence of matches from the input text. Regular expressions are a powerful tool to extract strings with certain characteristics. The pattern `\p{L}+` means that it will extract letters without accents (e.g., it will extract "Acuna" from "Acuña"). `setGaps=False` means that it will keep applying the rule until it can't extract new words. You have to set the input column from the incoming dataframe (in our case the `review` column) and the new column that will be added (e.g., `words`).

```
tokenizer = RegexTokenizer().setGaps(False)\
    .setPattern("\\p{L}+")\
    .setInputCol("review")\
    .setOutputCol("words")

review_words_df = tokenizer.transform(imdb_reviews_df)
print(review_words_df)
```

```
DataFrame[id: string, review: string, score: double, words:
array<string>]
```

Applying the transformation doesn't actually do anything until you apply an action. But as you can see, a new column `words` of type `array of string` was added by the transformation. We can see how it looks:

```
review_words_df.show(5)
```

```
+-----+-----+-----+-----+
|      id|      review|score|      words|
```



```
+-----+-----+-----+-----+
|pos_10006|In this "critical...| 1.0|[in, this, critic...|
|pos_10013|Like one of the p...| 1.0|[like, one, of, t...|
|pos_10022|Aro Tolbukhin bur...| 1.0|[aro, tolbukhin, ...|
|pos_10033|The movie Titanic...| 1.0|[the, movie, tita...|
| pos_1003|Another Aussie ma...| 1.0|[another, aussie,...|
+-----+-----+-----+-----+
only showing top 5 rows
```

Now, we want to match every word from `sentiment_df` in the array `words` shown before. One way of doing this is to *explode* the column `words` to create a row for each element in that list. Then, we would join that result with the dataframe `sentiment` to continue further.

```
review_words_df.select('id',
fn.explode('words').alias('word1')).show(5)
```

```
+-----+-----+
|      id|      word1|
+-----+-----+
|pos_10006|      in|
|pos_10006|      this|
|pos_10006|  critically|
|pos_10006|    acclaimed|
|pos_10006|psychological|
+-----+-----+
only showing top 5 rows
```

Now if we join that with `sentiment`, we can see if there are positive and negative words in each review:

```
review_word_sentiment_df = review_words_df.\
    select('id', fn.explode('words').alias('word')).\
    join(sentiments_df, 'word')
review_word_sentiment_df.show(5)
```

```
+-----+-----+-----+
|      word|      id|sentiment|
+-----+-----+-----+
|  acclaimed|pos_10006|      1|
|celebrated|pos_10006|      1|
|  troubling|pos_10006|     -1|
|   mystery|pos_10006|     -1|
|   deadly|pos_10006|     -1|
+-----+-----+-----+
only showing top 5 rows
```

Check the unique sentiment column values. Sentiments should be +1 and -1 for positive and negative word sentiments respectively.

```
review_word_sentiment_df.groupBy("sentiment").count().show()
```

```
+-----+-----+
|sentiment| count|
+-----+-----+
|         1|257274|
|        -1|241972|
+-----+-----+
```

Now we can simply average the sentiment per review id and, say, pick positive when the average is above 0, and negative otherwise.

```
simple_sentiment_prediction_df = review_word_sentiment_df.\
    groupBy('id').\
    agg(fn.avg('sentiment').alias('avg_sentiment')).\
    withColumn('predicted', fn.when(fn.col('avg_sentiment') > 0,\
1.0).otherwise(0.))
simple_sentiment_prediction_df.show(5)
```

```
+-----+-----+-----+
|      id| avg_sentiment|predicted|
+-----+-----+-----+
|pos_10149| 0.42857142857142855|      1.0|
|pos_10377| 0.5384615384615384|      1.0|
| pos_1299| 0.09090909090909091|      1.0|
| pos_2228| -0.14285714285714285|      0.0|
| pos_5052| 0.7777777777777778|      1.0|
+-----+-----+-----+
only showing top 5 rows
```

Now, lets compute the accuracy of our prediction

```
imdb_reviews_df.\
    join(simple_sentiment_prediction_df, 'id').\
    select(fn.expr('float(score == predicted)').alias('correct')).\
    select(fn.avg('correct')).\
    show()
```

```
+-----+
| avg(correct)|
+-----+
|0.732231471106131|
+-----+
```

## A data-driven sentiment prediction

First, we need to create a sequence to take from raw text to term frequency. This is necessary because we don't know the number of tokens in the text and therefore we need to *estimate* such quantity from the data.

```
# we obtain the stop words from a website
import requests
stop_words =
requests.get('http://ir.dcs.gla.ac.uk/resources/linguistic_utils/stop_
words').text.split()
stop_words[0:10]

['a',
 'about',
 'above',
 'across',
 'after',
 'afterwards',
 'again',
 'against',
 'all',
 'almost']

from pyspark.ml.feature import StopWordsRemover
sw_filter = StopWordsRemover()\
    .setStopWords(stop_words)\
    .setCaseSensitive(False)\
    .setInputCol("words")\
    .setOutputCol("filtered")
```

Count Vectorizer

```
from pyspark.ml.feature import CountVectorizerizer

# we will remove words that appear in 5 docs or less
cv = CountVectorizer(minTF=1., minDF=5., vocabSize=2**17)\
    .setInputCol("filtered")\
    .setOutputCol("tf")

# we now create a pipelined transformer
cv_pipeline = Pipeline(stages=[tokenizer, sw_filter,
cv]).fit(imdb_reviews_df)

# now we can make the transformation between the raw text and the
counts
cv_pipeline.transform(imdb_reviews_df).show(5)

+-----+-----+-----+-----+
+-----+-----+-----+-----+
```

id  filtered	review  tf	score	words
pos_10006	In this "critical...	1.0	[in, this, critic...
	[critically, accl... (26677,[0,1,3,4,5...		
pos_10013	Like one of the p...	1.0	[like, one, of, t... [like,
	previous, ... (26677,[1,2,3,4,5...		
pos_10022	Aro Tolbukhin bur...	1.0	[aro, tolbukhin, ... [aro,
	tolbukhin, ... (26677,[0,1,2,12,...		
pos_10033	The movie Titanic...	1.0	[the, movie, tita... [movie,
	titanic, ... (26677,[0,1,2,3,4...		
pos_1003	Another Aussie ma...	1.0	[another, aussie,... [aussie,
	masterpi... (26677,[4,5,9,24,...		

only showing top 5 rows

The term frequency vector is represented with a sparse vector. We have 26,677 terms.

```
len(cv_pipeline.stages[-1].vocabulary)
26677
```

Finally, we build another pipeline that takes the output of the previous pipeline and *lowers* the terms of documents that are very common.

```
from pyspark.ml.feature import IDF
idf = IDF().\
    setInputCol('tf').\
    setOutputCol('tfidf')

idf_pipeline = Pipeline(stages=[cv_pipeline,
idf]).fit(imdb_reviews_df)


idf_pipeline.transform(imdb_reviews_df).select('tfidf').show(1, False)
```

```
-----+
|tfidf
|
+-----
```

| (26677,  
[0,1,3,4,5,7,8,9,10,14,21,24,26,29,30,31,37,38,46,48,51,54,59,64,88,11  
6,125,129,150,189,200,202,222,227,230,232,239,243,255,261,263,269,328,  
347,356,364,368,377,392,432,451,454,472,486,514,517,521,540,567,570,58  
3,628,630,631,633,647,650,651,673,705,747,749,952,1058,1083,1158,1254,  
1303,1316,1320,1325,1399,1400,1464,1472,1474,1530,1574,1654,1693,1791,

1815,1817,1826,1845,1917,1954,2008,2009,2011,2022,2117,2188,2272,2453,  
2464,2882,2946,3018,3109,3289,3309,3538,3622,3624,3642,3794,3922,4043,  
4354,4513,4815,5050,5158,5186,5576,5745,5947,6089,6189,6263,6357,6360,  
7205,7219,7508,7972,9998,10080,10103,10347,10535,10684,11840,12761,136  
33,14591,15021,15624,16078,18058,20532,22354,24267,24420,25217],  
[6.399789021643247,1.5355439272599614,1.751599547551507,0.502897453316  
8351,0.7612090942068761,0.9527872680847915,1.051577880482202,2.3754916  
62124171,1.1869590489601298,1.3491156498373038,1.5198200397035864,1.68  
87375926197539,1.53920585864879,1.6292725389623013,1.6586681558248935,  
1.6538478593700623,1.8765733504775144,3.705648340531702,1.924736751856  
1371,3.905063519241778,4.097206442469843,1.9536597913253875,2.04457782  
4871977,2.2186516584178886,2.4178115020031847,2.4609176711886604,2.553  
139840304409,5.076694851430274,5.339278027855957,2.950416571033215,2.8  
99008604321704,2.8228278777726548,2.8577509748566374,2.912166688899215  
5,2.8732007203658316,3.0349289879552064,2.895384098722744,3.1456653623  
286295,2.951946793213983,8.994525472110968,3.2555798084398138,2.949652  
337140459,3.125425480980883,3.210947654419045,3.2059995988016756,3.257  
6566523846524,3.216917821405549,3.2350452059981056,3.3365738675364547,  
3.437071833871393,3.3980424899656576,3.360479388390009,3.4802805882026  
293,3.522727278449887,3.5755907680069545,3.5336265689079225,3.44707191  
72059763,3.6969516250112218,3.5784520002879865,3.5755907680069545,3.57  
2737699024548,3.7987343193211642,3.6223829295137135,3.734753989657757,  
3.62538143250997,4.135206555942377,3.7083061671141477,3.66520292669664  
15,3.734753989657757,3.8462752640901643,3.916071026025706,4.2050926834  
06543,4.1105139433520055,4.115403928646197,4.205092683406543,4.2240377  
69648993,4.367769329173078,9.116653198578526,4.3306133522849874,4.3836  
67915240877,23.37816324742329,4.4096434016441375,4.577595018155139,4.4  
95459321228993,4.499049989359722,4.467188887290738,4.467188887290738,4  
.506270237333209,4.546941277064137,4.662839298024749,5.206690177222234  
,9.93523160767166,4.728508401532606,4.692949099496119,4.79879493426077  
8,4.779563572332891,4.7421760402612705,10.128152140046785,5.0962331816  
57924,4.756033074922696,4.89022914022041,5.129458829286244,4.933714252  
160149,4.900924429337158,4.996772388127286,5.192197169919667,5.2668586  
98688687,5.185028680441055,5.5833763207803555,5.251473779849207,5.3391  
793602683135,5.330880557453618,5.863991226009044,5.63803473331822,5.70  
7830495253761,5.472710752892835,5.5115505862091,5.501698289766088,12.1  
67239670431618,12.510940184284935,5.672323806796851,6.10131941231521,5  
.795937762764028,6.276523501340301,5.921978483659393,5.967788019690687  
,6.06622809250394,12.64001722656008,6.083619835215809,6.06622809250394  
,39.258912987565495,6.10131941231521,6.137687056486085,6.2980297065612  
64,6.298029706561264,6.32000861328004,13.66166847409206,6.725473721388  
203,6.794466592875155,13.98235377424242,14.071257299384087,6.830834237  
04603,7.035628649692043,6.99117688712121,7.082148665326936,7.182232123  
883919,7.293457758994143,7.354082380810578,7.4876137734351005,7.561721  
745588822,7.824086010056313,8.047229561370523,8.047229561370523,8.1807  
60953995046,8.334911633822305,8.334911633822305])|

+-----  
-----  
-----





only showing top 1 row

Therefore, the `idf_pipeline` takes the raw text from the dataframe `imdb_reviews_df` and creates a feature vector called `tfidf`!

```
tfidf_df = idf_pipeline.transform(imdb_reviews_df)
print("tfidf_df shape: ", tfidf_df.count(), len(tfidf_df.columns))
```

```
tfidf_df shape: 25000 7
```

```
tfidf_df.show(5)
```

```
+-----+-----+-----+-----+
+-----+-----+-----+-----+
|      id|      review|score|      words|
|filtered|          tf|      tfidf|
+-----+-----+-----+-----+
+-----+-----+-----+-----+
|pos_10006|In this "critical...| 1.0|[in, this, critic...|
|critically, accl...|(26677,[0,1,3,4,5...|(26677,[0,1,3,4,5...|
|pos_10013|Like one of the p...| 1.0|[like, one, of, t...|[like,
previous, ...|(26677,[1,2,3,4,5...|(26677,[1,2,3,4,5...|
|pos_10022|Aro Tolbukhin bur...| 1.0|[aro, tolbukhin, ...|[aro,
tolbukhin, ...|(26677,[0,1,2,12,...|(26677,[0,1,2,12,...|
|pos_10033|The movie Titanic...| 1.0|[the, movie, tita...|[movie,
titanic, ...|(26677,[0,1,2,3,4...|(26677,[0,1,2,3,4...|
| pos_1003|Another Aussie ma...| 1.0|[another, aussie,...|[aussie,
masterpi...|(26677,[4,5,9,24,...|(26677,[4,5,9,24,...|
+-----+-----+-----+-----+
+-----+-----+-----+-----+
```

only showing top 5 rows

The cell below prints out the 'tf' and 'tfidf' columns for the first 10 rows of the `tfidf_df`. Note that the `tfidf` column is transformed to account for the frequency with which the word appears in the corpus. Words that appear more often are penalized more than words that do not appear as frequently in the corpus.

```
tfidf_df.limit(10).toPandas().loc[:10, ['tf', 'tfidf']]
```

```
tf \
0 (12.0, 5.0, 0.0, 3.0, 1.0, 1.0, 0.0, 1.0, 1.0,...
1 (0.0, 2.0, 4.0, 1.0, 1.0, 3.0, 3.0, 0.0, 0.0, ...
2 (6.0, 2.0, 1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...
3 (6.0, 1.0, 3.0, 7.0, 3.0, 2.0, 1.0, 0.0, 2.0, ...
4 (0.0, 0.0, 0.0, 0.0, 1.0, 1.0, 0.0, 0.0, 0.0, ...
5 (14.0, 7.0, 1.0, 6.0, 2.0, 0.0, 2.0, 3.0, 1.0,...
6 (6.0, 2.0, 3.0, 1.0, 2.0, 1.0, 0.0, 0.0, 1.0, ...
7 (6.0, 0.0, 0.0, 2.0, 0.0, 1.0, 0.0, 0.0, 0.0, ...
```

```

8 (4.0, 0.0, 0.0, 4.0, 0.0, 1.0, 0.0, 0.0, 0.0, ...
9 (6.0, 1.0, 0.0, 2.0, 1.0, 1.0, 1.0, 1.0, 0.0, ...

                                tfidf
0 (6.399789021643247, 1.5355439272599614, 0.0, 1...
1 (0.0, 0.6142175709039845, 1.9611159574304948, ...
2 (3.1998945108216237, 0.6142175709039845, 0.490...
3 (3.1998945108216237, 0.30710878545199227, 1.47...
4 (0.0, 0.0, 0.0, 0.0, 0.5028974533168351, 0.761...
5 (7.466420525250455, 2.149761498163946, 0.49027...
6 (3.1998945108216237, 0.6142175709039845, 1.470...
7 (3.1998945108216237, 0.0, 0.0, 1.1677330317010...
8 (2.1332630072144156, 0.0, 0.0, 2.3354660634020...
9 (3.1998945108216237, 0.30710878545199227, 0.0,...

```

## Data science pipeline for estimating sentiments

First, let's split the data into training, validation, and testing.

```

training_df, validation_df, testing_df =
imdb_reviews_df.randomSplit([0.6, 0.3, 0.1], seed=0)

training_df.limit(5).toPandas()

```

	id	review	score
0	neg_1	Robert DeNiro plays the most unbelievably inte...	0.0
1	neg_10	This film had a lot of promise, and the plot w...	0.0
2	neg_100	OK its not the best film I've ever seen but at...	0.0
3	neg_10000	Airport '77 starts as a brand new luxury 747 p...	0.0
4	neg_10001	This film lacked something I couldn't put my f...	0.0

```

[training_df.count(), validation_df.count(), testing_df.count()]

[14962, 7531, 2507]

```

Logistic Regression

```

from pyspark.ml.classification import LogisticRegression

lr = LogisticRegression().\
    setLabelCol('score').\
    setFeaturesCol('tfidf').\
    setRegParam(0.0).\
    setMaxIter(100).\
    setElasticNetParam(0.)

```

Lets create a pipeline transformation by chaining the `idf_pipeline` with the logistic regression step (`lr`)

```
lr_pipeline = Pipeline(stages=[idf_pipeline, lr]).fit(training_df)
```

The next cell defines a class capable of calculating ROC and PR curves. The following cell creates a ROC curve using transformed data from the pipeline above.

```
# see https://stackoverflow.com/questions/52847408/pyspark-extract-roc-curve
from pyspark.mllib.evaluation import BinaryClassificationMetrics

# Scala version implements .roc() and .pr()
# Python:
https://spark.apache.org/docs/latest/api/python/_modules/pyspark/mllib/common.html
# Scala:
https://spark.apache.org/docs/latest/api/java/org/apache/spark/mllib/evaluation/BinaryClassificationMetrics.html
class CurveMetrics(BinaryClassificationMetrics):
    def __init__(self, *args):
        super(CurveMetrics, self).__init__(*args)

    def _to_list(self, rdd):
        points = []

        for row in rdd.collect():
            points += [(float(row._1()), float(row._2()))]
        return points

    def get_curve(self, method):
        rdd = getattr(self._java_model, method)().toJavaRDD()
        return self._to_list(rdd)

import matplotlib.pyplot as plt


# Create a Pipeline estimator and fit on train DF, predict on test DF
predictions = lr_pipeline.transform(validation_df)

# Returns as a list (false positive rate, true positive rate)
preds = predictions.select('score', 'probability').rdd.map(lambda row:
(float(row['probability'][1]), float(row['score'])))
points = CurveMetrics(preds).get_curve('roc')

plt.figure()
x_val = [x[0] for x in points]
y_val = [x[1] for x in points]
plt.title('IMDB Sentiment Analysis ROC Curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.plot(x_val, y_val)
```

```
/usr/local/lib/python3.9/dist-packages/pyspark/sql/context.py:157:
FutureWarning: Deprecated in 3.0.0. Use
SparkSession.builder.getOrCreate() instead.
  warnings.warn(
```

```
[<matplotlib.lines.Line2D at 0x7fed7c131610>]
```



Lets estimate the accuracy:

```
lr_pipeline.transform(validation_df).\
  select(fn.expr('float(prediction = score)').alias('correct')).\
  select(fn.avg('correct')).show()
```

```
+-----+
|      avg(correct) |
+-----+
|0.8360111538972248|
+-----+
```

The performance is much better than before.

The problem however is that we are overfitting because we have many features compared to the training examples:

For example, if we look at the weights of the features, there is a lot of noise:

```
import pandas as pd
vocabulary = idf_pipeline.stages[0].stages[-1].vocabulary
weights = lr_pipeline.stages[-1].coefficients.toArray()
print("num weights:", len(weights))
print("num rows:", validation_df.count())

coeffs_df = pd.DataFrame({'word': vocabulary, 'weight': weights})
coeffs_df.head()
```

num weights: 26677  
num rows: 7531

	word	weight
0	br	-0.105915
1	s	0.181536
2	movie	-0.336467
3	film	0.194032
4	t	-0.783280

The most negative words are:

```
coeffs_df.sort_values('weight').head(5)
```

	word	weight
23304	zoolander	-6.642286
21131	grossness	-6.096098
26293	residue	-5.677060
23883	publicize	-5.493201
23282	dodgeball	-5.350357

And the most positive:

```
coeffs_df.sort_values('weight', ascending=False).head(5)
```

	word	weight
20792	choco	7.881534
26406	silhouetted	7.406967
26318	eliciting	6.151619
21679	praiseworthy	6.142447
13496	shrewd	5.979843

But none of them make sense. What is happening? We are overfitting the data. Those words that don't make sense are capturing just noise in the reviews.

## Spark allows us to fit elastic net regularization easily

```
lambda_par = 0.02
alpha_par = 0.3
en_lr = LogisticRegression().\
    setLabelCol('score').\
    setFeaturesCol('tfidf').\
    setRegParam(lambda_par).\
    setMaxIter(100).\
    setElasticNetParam(alpha_par)
```

And we define a new Pipeline with all steps combined

```
en_lr_estimator = Pipeline(stages=[tokenizer, sw_filter, cv, idf,
en_lr])

en_lr_pipeline = en_lr_estimator.fit(training_df)
```

Let's look at the performance

```
en_lr_pipeline.transform(validation_df).select(fn.avg(fn.expr('float(p
rediction = score)'))).show()
```

```
+-----+
|avg((prediction = score))|
+-----+
|          0.8688089231177798|
+-----+
```

We improve performance slightly, but whats more important is that we improve the understanding of the word sentiments. Lets look at the weights:

```
en_weights = en_lr_pipeline.stages[-1].coefficients.toArray()
en_coefs_df = pd.DataFrame({'word':
en_lr_pipeline.stages[2].vocabulary, 'weight': en_weights})
```

The most negative words all make sense ("worst" is *actually* more negative than than "worse")!

```
en_coefs_df.sort_values('weight').head(15)
```

	word	weight
105	worst	-0.369188
262	waste	-0.334131
190	awful	-0.259531
12	bad	-0.238970
606	poorly	-0.185563
522	dull	-0.185538
1109	disappointment	-0.181951

184	boring	-0.180375
179	poor	-0.177850
243	worse	-0.173731
351	horrible	-0.172044
1376	redeeming	-0.163199
1180	disappointing	-0.154673
579	avoid	-0.152653
1032	laughable	-0.151916

Same thing with positive words

```
en_coefs_df.sort_values('weight', ascending=False).head(15)
```

	word	weight
13	great	0.283820
161	excellent	0.240269
221	wonderful	0.200133
26	best	0.168678
300	favorite	0.161667
227	perfect	0.157898
286	amazing	0.142479
889	incredible	0.139636
270	loved	0.133010
2047	refreshing	0.129359
1962	captures	0.127790
311	enjoyed	0.125049
700	perfectly	0.123770
309	today	0.122013
3047	flawless	0.120523

Are there words with *literally* zero importance for predicting sentiment? Yes, and most of them!

```
en_coefs_df.query('weight == 0.0').shape
(19609, 2)
```

In fact, approximately 95% of features are not needed to achieve a **better** performance than all previous models!

```
en_coefs_df.query('weight == 0.0').shape[0]/en_coefs_df.shape[0]
0.9502786527744124
```

Let's look at these *neutral* words

```
en_coefs_df.query('weight == 0.0').head(15)
```

	word	weight
0	br	0.0

3	film	0.0
5	like	0.0
9	story	0.0
10	really	0.0
11	people	0.0
14	don	0.0
15	way	0.0
17	movies	0.0
18	think	0.0
19	characters	0.0
20	character	0.0
29	little	0.0
31	know	0.0
32	man	0.0

But, did we choose the right  $\lambda$  and  $\alpha$  parameters? We should run an experiment where we try different combinations of them. Fortunately, Spark let us do this by using a grid - a method that generates combination of parameters.

```
from pyspark.ml.tuning import ParamGridBuilder
```

We need to build a new estimator pipeline

```
en_lr_estimator.getStages()

[RegexTokenizer_18e623850b1a,
 StopWordsRemover_49bcf0e3b650,
 CountVectorizer_789893e3001b,
 IDF_507b4bb23612,
 LogisticRegression_b427af0d8b7f]

grid = ParamGridBuilder().\
    addGrid(en_lr.regParam, [0., 0.01, 0.02]).\
    addGrid(en_lr.elasticNetParam, [0., 0.2, 0.4]).\
    build()
```

This is the list of parameters that we will try:

```
grid

[Param(parent='LogisticRegression_b427af0d8b7f', name='regParam',
 doc='regularization parameter (>= 0).'): 0.0,
 Param(parent='LogisticRegression_b427af0d8b7f',
 name='elasticNetParam', doc='the ElasticNet mixing parameter, in range
 [0, 1]. For alpha = 0, the penalty is an L2 penalty. For alpha = 1, it
 is an L1 penalty.'): 0.0},
 {Param(parent='LogisticRegression_b427af0d8b7f', name='regParam',
 doc='regularization parameter (>= 0).'): 0.0,
 Param(parent='LogisticRegression_b427af0d8b7f',
```



```

name='elasticNetParam', doc='the ElasticNet mixing parameter, in range
[0, 1]. For alpha = 0, the penalty is an L2 penalty. For alpha = 1, it
is an L1 penalty.'): 0.2},
    {Param(parent='LogisticRegression_b427af0d8b7f', name='regParam',
doc='regularization parameter ( $\geq 0$ ).'): 0.0,
    Param(parent='LogisticRegression_b427af0d8b7f',
name='elasticNetParam', doc='the ElasticNet mixing parameter, in range
[0, 1]. For alpha = 0, the penalty is an L2 penalty. For alpha = 1, it
is an L1 penalty.'): 0.4},
    {Param(parent='LogisticRegression_b427af0d8b7f', name='regParam',
doc='regularization parameter ( $\geq 0$ ).'): 0.01,
    Param(parent='LogisticRegression_b427af0d8b7f',
name='elasticNetParam', doc='the ElasticNet mixing parameter, in range
[0, 1]. For alpha = 0, the penalty is an L2 penalty. For alpha = 1, it
is an L1 penalty.'): 0.0},
    {Param(parent='LogisticRegression_b427af0d8b7f', name='regParam',
doc='regularization parameter ( $\geq 0$ ).'): 0.01,
    Param(parent='LogisticRegression_b427af0d8b7f',
name='elasticNetParam', doc='the ElasticNet mixing parameter, in range
[0, 1]. For alpha = 0, the penalty is an L2 penalty. For alpha = 1, it
is an L1 penalty.'): 0.2},
    {Param(parent='LogisticRegression_b427af0d8b7f', name='regParam',
doc='regularization parameter ( $\geq 0$ ).'): 0.01,
    Param(parent='LogisticRegression_b427af0d8b7f',
name='elasticNetParam', doc='the ElasticNet mixing parameter, in range
[0, 1]. For alpha = 0, the penalty is an L2 penalty. For alpha = 1, it
is an L1 penalty.'): 0.4},
    {Param(parent='LogisticRegression_b427af0d8b7f', name='regParam',
doc='regularization parameter ( $\geq 0$ ).'): 0.02,
    Param(parent='LogisticRegression_b427af0d8b7f',
name='elasticNetParam', doc='the ElasticNet mixing parameter, in range
[0, 1]. For alpha = 0, the penalty is an L2 penalty. For alpha = 1, it
is an L1 penalty.'): 0.0},
    {Param(parent='LogisticRegression_b427af0d8b7f', name='regParam',
doc='regularization parameter ( $\geq 0$ ).'): 0.02,
    Param(parent='LogisticRegression_b427af0d8b7f',
name='elasticNetParam', doc='the ElasticNet mixing parameter, in range
[0, 1]. For alpha = 0, the penalty is an L2 penalty. For alpha = 1, it
is an L1 penalty.'): 0.2},
    {Param(parent='LogisticRegression_b427af0d8b7f', name='regParam',
doc='regularization parameter ( $\geq 0$ ).'): 0.02,
    Param(parent='LogisticRegression_b427af0d8b7f',
name='elasticNetParam', doc='the ElasticNet mixing parameter, in range
[0, 1]. For alpha = 0, the penalty is an L2 penalty. For alpha = 1, it
is an L1 penalty.'): 0.4}]

all_models = []
for j in range(len(grid)):
    print("Fitting model {}".format(j+1))

```

```

    model = en_lr_estimator.fit(training_df, grid[j])
    all_models.append(model)

```

```

Fitting model 1
Fitting model 2
Fitting model 3
Fitting model 4
Fitting model 5
Fitting model 6
Fitting model 7
Fitting model 8
Fitting model 9

```

```

# estimate the accuracy of each of them:
accuracies = [m.\
    transform(validation_df).\
    select(fn.avg(fn.expr('float(score =\
prediction)'))).alias('accuracy')).\
    first().\
    accuracy for m in all_models]

```

```

accuracies

```

```

[0.8389324126941973,
 0.8389324126941973,
 0.8389324126941973,
 0.8633647589961493,
 0.8794316823794981,
 0.8729252423316957,
 0.8678794316823795,
 0.8761120701102111,
 0.8600451467268623]

```

```

import numpy as np

best_model_idx = np.argmax(accuracies)
print("best model index =", best_model_idx)

```

```

best model index = 4

```

So the best model we found has the following parameters

```

grid[best_model_idx]

```

```

{Param(parent='LogisticRegression_b427af0d8b7f', name='regParam',
doc='regularization parameter (>= 0).'): 0.01,
 Param(parent='LogisticRegression_b427af0d8b7f',
name='elasticNetParam', doc='the ElasticNet mixing parameter, in range
[0, 1]. For alpha = 0, the penalty is an L2 penalty. For alpha = 1, it
is an L1 penalty.'): 0.2}

```

```
best_model = all_models[best_model_idx]
accuracies[best_model_idx]
0.8794316823794981

# estimate generalization performance
best_model.\
  transform(testing_df).\
  select(fn.avg(fn.expr('float(score =
prediction)'))).alias('accuracy')).\
  show()
```

accuracy
0.8687674511368169

## Finally, predicting tweet sentiments

Now we can use this model to predict sentiments on Twitter

```
tweets_df = spark.read.parquet('tweets.parquet')
```

We have 1K tweets from each candidate

```
tweets_df.groupby('handle').agg(fn.count('*')).show()
```

handle	count(1)
@HillaryClinton	1000
@realDonaldTrump	1000

We can now predict the sentiment of the Tweet using our best model, we need to rename the column so that it matches our previous pipeline (`review => ...`)

```
best_model.transform(tweets_df.withColumnRenamed('text',
'review')).select('handle', 'review', 'prediction').show()
```

handle	review	prediction
@HillaryClinton	RT @ZekeJMiller: ...	1.0
@HillaryClinton	"She's just out t...	1.0

@HillaryClinton	We're going to ma...	0.0
@HillaryClinton	Don't boo. Vote! ...	0.0
@HillaryClinton	This Republican d...	0.0
@HillaryClinton	Hillary teamed up...	1.0
@HillaryClinton	RT @mayaharris_: ...	1.0
@HillaryClinton	"It was overwhelm...	1.0
@HillaryClinton	Great step forward...	1.0
@HillaryClinton	"I feel like I'm ...	1.0
@HillaryClinton	Nobody here was "...	1.0
@HillaryClinton	For those few peo...	0.0
@HillaryClinton	Remember, don't b...	1.0
@HillaryClinton	Too many talented...	1.0
@HillaryClinton	There are hundred...	0.0
@HillaryClinton	It's 3:20am. As g...	1.0
@HillaryClinton	Trump stood on a ...	0.0
@HillaryClinton	Donald Trump said...	1.0
@HillaryClinton	RT @timkaine: 39 ...	1.0
@HillaryClinton	Trump wants to br...	1.0

+-----+-----+-----+

only showing top 20 rows

Now, lets summarize our results in a graph!


```
%matplotlib inline
import seaborn

sentiment_pd = best_model.\
    transform(tweets_df.withColumnRenamed('text', 'review')).\
    groupby('handle').\
    agg(fn.avg('prediction').alias('ave_prediction'),
(2*fn.stddev('prediction')/fn.sqrt(fn.count('*'))).alias('std_err')).\
    toPandas()

sentiment_pd.head()

   handle  ave_prediction  std_err
0  @HillaryClinton      0.571  0.031318
1  @realDonaldTrump      0.702  0.028942

sentiment_pd.plot(x='handle', y='ave_prediction', xerr='std_err',
kind='barh');
```



But let's examine some "negative" tweets by Trump

```
best_model.\
  transform(tweets_df.withColumnRenamed('text', 'review')).\
  where(fn.col('handle') == '@realDonaldTrump').\
  where(fn.col('prediction') == 0).\
  select('review').\
  take(5)
```

```
[Row(review='Moderator: Hillary paid $225,000 by a Brazilian bank for
a speech that called for "open borders." That's a quote! #Debate
#BigLeagueTruth'),
 Row(review='TRUMP & CLINTON ON IMMIGRATION\n#Debate
#BigLeagueTruth https://t.co/OP4c7Jc8Ad'),
 Row(review='Hillary is too weak to lead on border security-no
solutions, no ideas, no credibility.She supported NAFTA, worst deal in
US history. #Debate'),
 Row(review='One of my first acts as President will be to deport the
drug lords and then secure the border. #Debate #MAGA'),
 Row(review='Hillary Clinton will use American tax dollars to provide
amnesty for thousands of illegals. I will put...
https://t.co/ZpV33TfbR6')]
```

And Clinton

```
best_model.\
  transform(tweets_df.withColumnRenamed('text', 'review')).\
  where(fn.col('handle') == '@HillaryClinton').\
  where(fn.col('prediction') == 0).\
  select('review').\
  take(5)
```

```
[Row(review="We're going to make college debt-free for everyone in
America. See how much you could save with Hillary's plan at...
https://t.co/Fhzkubhpj7"),
 Row(review="Don't boo. Vote! https://t.co/tTgeqy51PU
https://t.co/9un3FUVxoG"),
 Row(review='This Republican dad is struggling with the idea of his
daughter growing up in a country led by Donald Trump.
https://t.co/Tn3rQqJJKp'),
 Row(review="For those few people knocking public service, hope you'll
reconsider answering the call to help others. Because we're stronger
together."),
 Row(review="There are hundreds of thousands more @AmeriCorps
applications than spots. Horrible! Let's expand it from 75,000 annual
members to 250,000.")]
```

```
from pyspark.ml import feature
```

```
from pyspark.sql import types
```

```
def probability_positive(probability_column):
    return float(probability_column[1])
func_probability_positive = fn.udf(probability_positive,
types.DoubleType())
```

```
prediction_probability_df = best_model.transform(validation_df).\
  withColumn('probability_positive',
func_probability_positive('probability')).\
  select('id', 'review', 'score', 'probability_positive')
prediction_probability_df.show()
```

id	review	score	probability_positive
neg_0	Story of a man wh...	0.0	0.4858256123045229
neg_1000	The plot for Desc...	0.0	0.23773092718405964
neg_10003	When I was little...	0.0	6.864567527598009E-4
neg_10006	I don't know who ...	0.0	0.08939747035805656
neg_10012	This movie must b...	0.0	0.00807233188176526
neg_10014	I saw this movie ...	0.0	0.45210546440878807
neg_10019	Kareena Kapoor in...	0.0	4.943247978200782E-4
neg_10022	Summer season is ...	0.0	0.6559859138052624
neg_10023	Shame on Yash Raj...	0.0	2.732321286786909...
neg_10024	First lesson that...	0.0	0.03837454919442762

neg_10026	I had some expect...	0.0	0.567648231964684
neg_1003	OK, I am not Japa...	0.0	0.16001109982402284
neg_10033	I was very disple...	0.0	0.3506335602422507
neg_10034	If there is one f...	0.0	2.447261518190302...
neg_10036	Sometime I fail t...	0.0	0.5942199783219011
neg_10041	The sight of Kare...	0.0	4.983562194804669E-6
neg_10050	A huge disappoint...	0.0	0.005242299674230...
neg_10052	Warner Bros. made...	0.0	0.20662075785815548
neg_10055	I grew up on the ...	0.0	0.3142833171137729
neg_10058	I was fascinated ...	0.0	0.12538714032358522

+-----+-----+-----+-----+

only showing top 20 rows